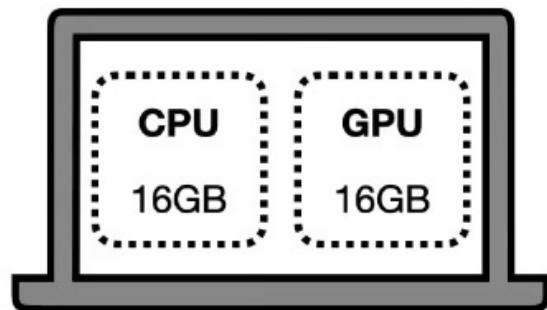


QLoRA

LLM fine-tuning made accessible

The Problem

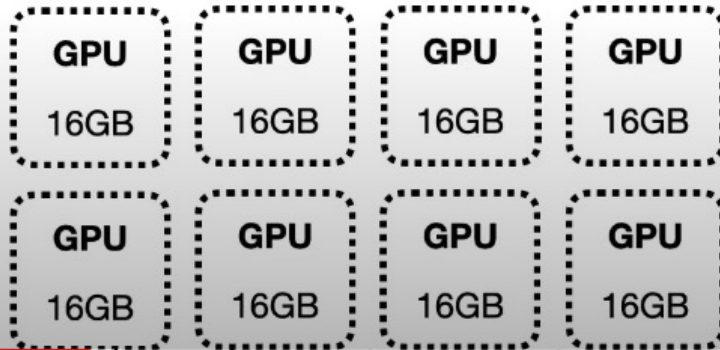
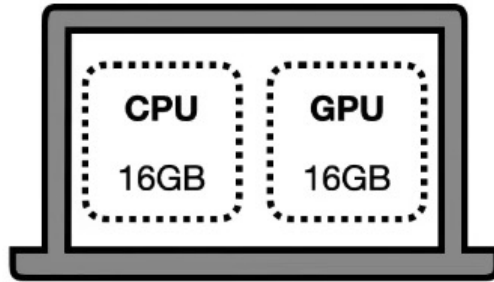
LLMs are (computationally) expensive



10B Parameter Model

The Problem

LLMs are (computationally) expensive



10B Parameter Model = 160GB!

Parameters (FP16) 20GB

Gradients (FP16) 20GB

Optimizer States
(FP32)

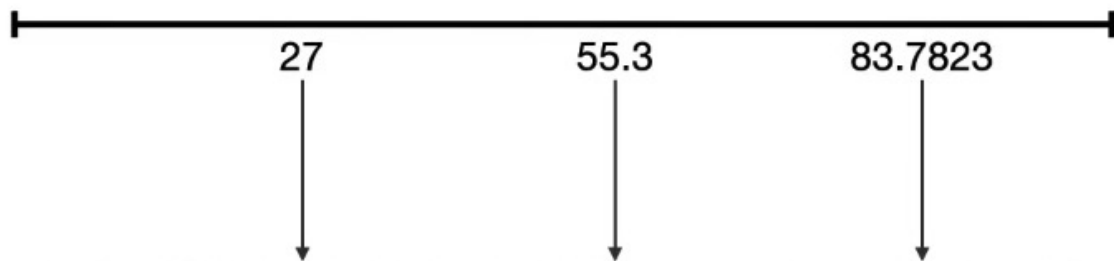
Momentum
Variance

120GB

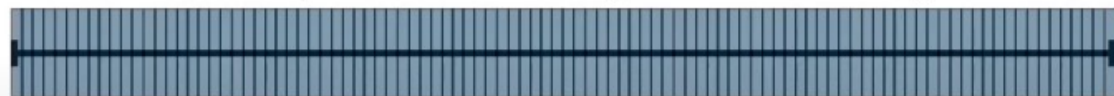
What is Quantization?

Quantization = splitting range into buckets

Any number
between 0 and 100



Quantized by
whole numbers

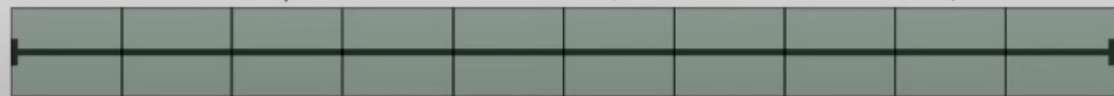


27

55

83

Quantized by 10s



20

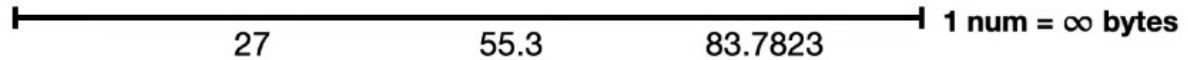
50

80

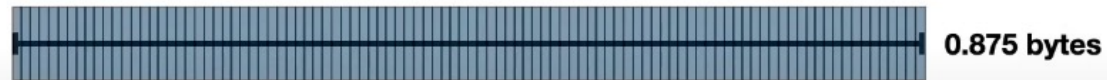
What is Quantization?

Quantization = splitting range into buckets

Any number
between 0 and 100

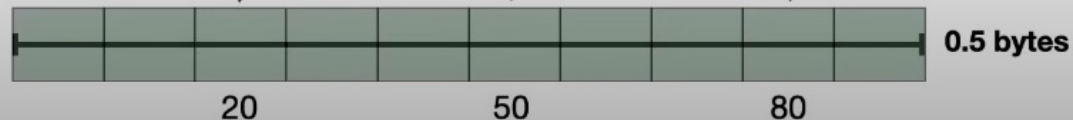


Quantized by
whole numbers



27 55 83

Quantized by 10s



Sign **Exponent** **Fraction / Mantissa**

32bit



0 10000001 11100100010101111100000 = 7.567856

16bit



0 10001 1110010001 = 7.566

```
tensor([[ 0.0031, -0.0438,  0.0494, ..., -0.0046, -0.0410,  0.0436],
        [-0.1013,  0.0394,  0.0787, ...,  0.0986,  0.0595,  0.0162],
        [-0.0859, -0.1227, -0.1209, ...,  0.1158,  0.0186, -0.0530],
        ...,
        [ 0.0804,  0.0725,  0.0638, ..., -0.0487, -0.0524, -0.1076],
        [-0.0200, -0.0406,  0.0663, ...,  0.0123,  0.0551, -0.0121],
        [-0.0041,  0.0865, -0.0013, ..., -0.0427, -0.0764,  0.1189]],
        dtype=torch.float16)
```



```
tensor([[  3, -47,  54, ...,  -5, -44,  47],
        [-104,  40,  81, ..., 101,  61,  17],
        [-89, -127, -125, ..., 120,  19, -55],
        ...,
        [ 82,  74,  65, ..., -49, -53, -109],
        [-21, -42,  68, ...,  13,  57, -12],
        [-4,  88, -1, ..., -43, -78, 121]],
        device='cuda:0', dtype=torch.int8, requires_grad=True)
```

8bit-Quantization

Ingredient 1: 4-bit NormalFloat

A better way to bucket numbers

4-bit e.g. 0101

$\Rightarrow 2^4 = 16$ unique combinations

$\Rightarrow 16$ buckets for quantizations

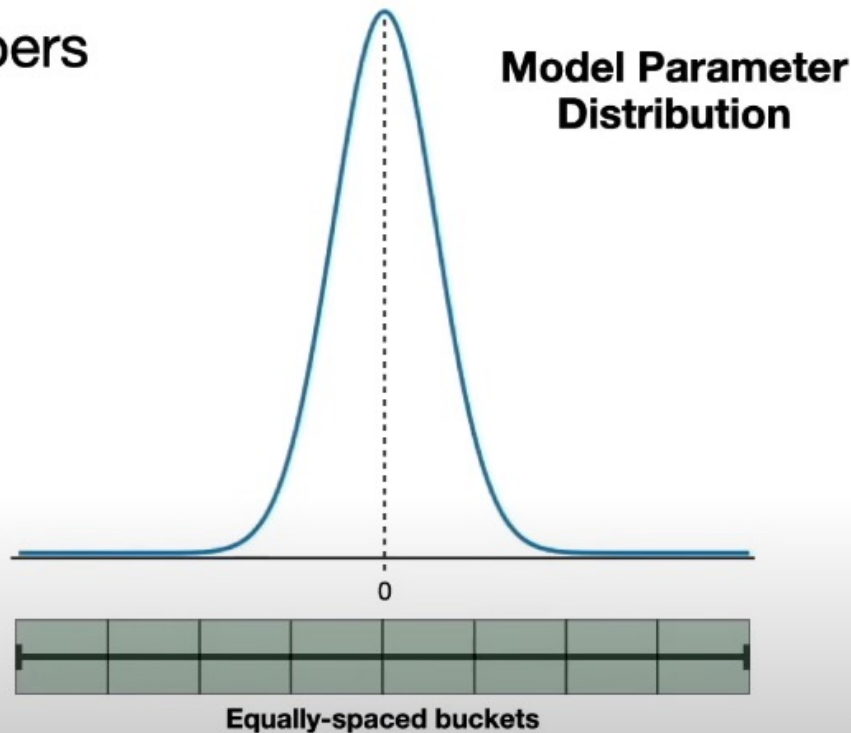
Ingredient 1: 4-bit NormalFloat

A better way to bucket numbers

4-bit e.g. 0101

$\Rightarrow 2^4 = 16$ unique combinations

$\Rightarrow 16$ buckets for quantizations



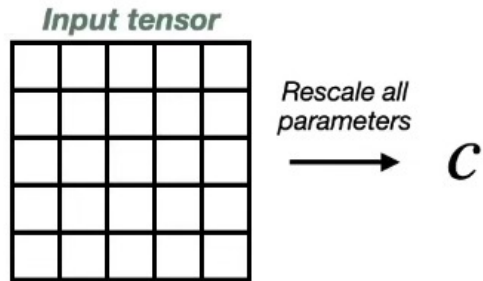
Ingredient 2: Double Quantization

Quantizing the Quantization Constants

$$x^{\text{Int8}} = \text{round} \left(\frac{127}{\text{absmax}(x^{\text{FP32}})} x^{\text{FP32}} \right)$$

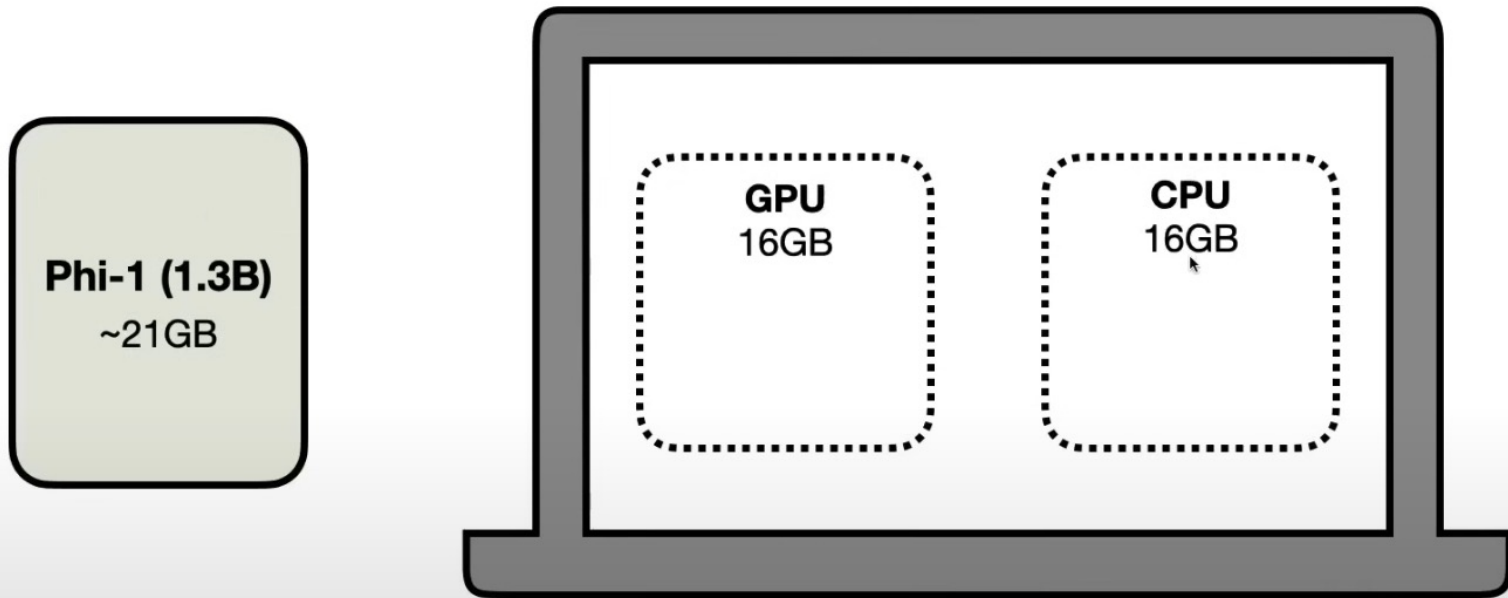
$$= \text{round} \left(c^{\text{FP32}} \cdot x^{\text{FP32}} \right)$$

↑
Takes up precious
memory



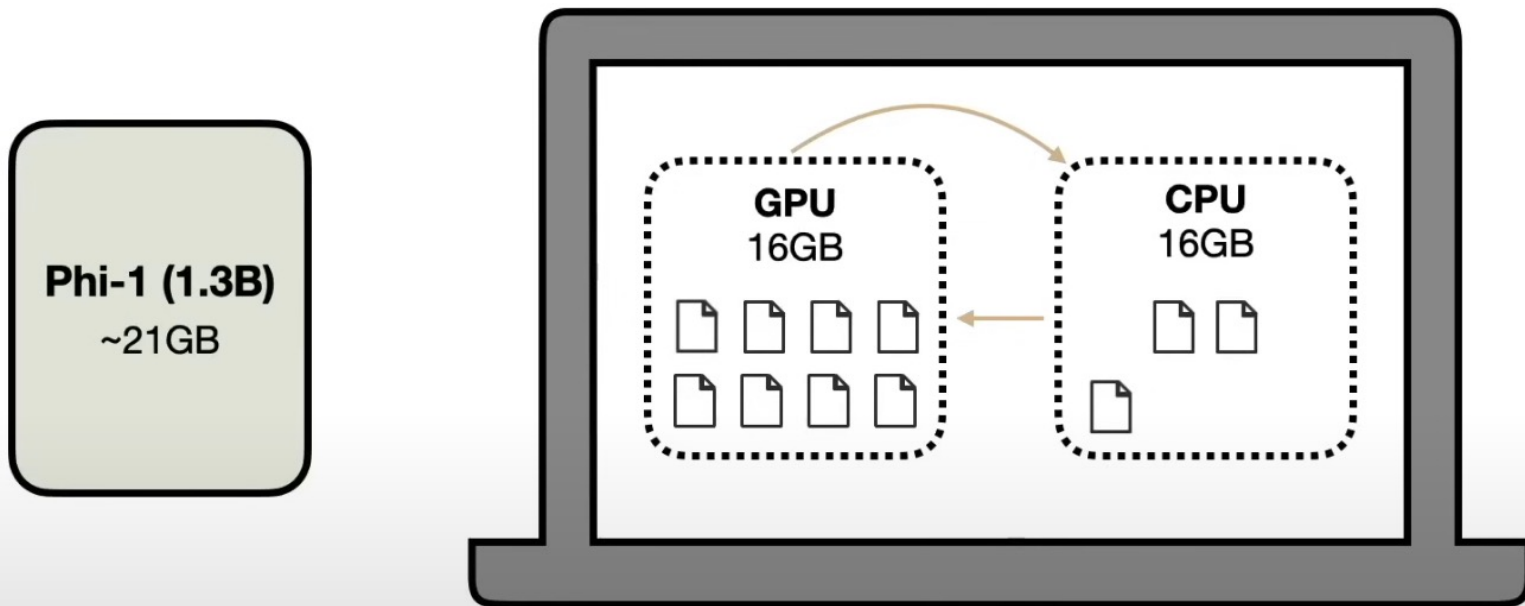
Ingredient 3: Paged Optimizer

Looping in your CPU



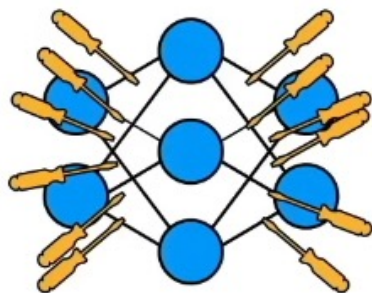
Ingredient 3: Paged Optimizer

Looping in your CPU

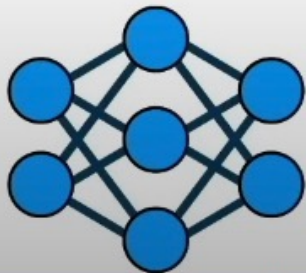


Ingredient 4: LoRA

Fine-tunes model by adding **small set** of trainable parameters



$x \quad h(x) \quad y$

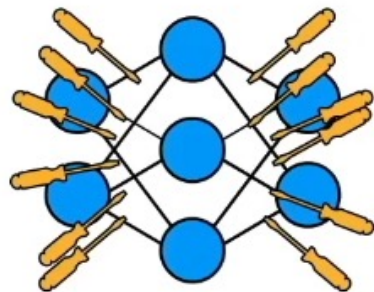


Full Fine-tuning: $h(x) = W_0 x$

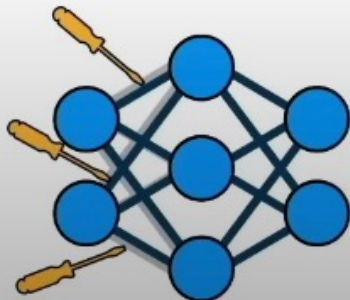
$$\begin{array}{c} \boxed{W_0} x = h(x) \\ \text{Trainable} \end{array}$$

Ingredient 4: LoRA

Fine-tunes model by adding **small set** of trainable parameters



$x \quad h(x) \quad y$



Full Fine-tuning: $h(x) = W_0x$

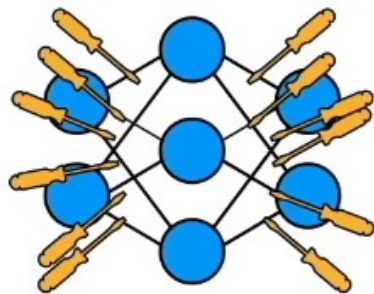
$$\begin{array}{|c|} \hline W_0 \\ \hline \end{array} x = \begin{array}{|c|} \hline h(x) \\ \hline \end{array}$$

Trainable

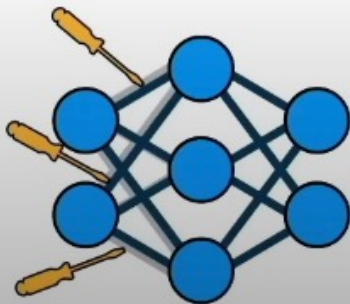
LoRA: $h(x) = W_0x + \Delta Wx = W_0x + BAx$

Ingredient 4: LoRA

Fine-tunes model by adding **small set** of trainable parameters



$x \quad h(x) \quad y$



Full Fine-tuning: $h(x) = W_0x$

$$\underbrace{\boxed{W_0}}_{\text{Trainable}} x = h(x)$$

LoRA: $h(x) = W_0x + \Delta Wx = W_0x + BAx$

$$\left(\underbrace{\boxed{W_0}}_{\text{Frozen}} + B \underbrace{A}_{\text{Trainable}} \right) x = h(x)$$

$$B \in \mathbb{R}^{d \times r}$$

$$A \in \mathbb{R}^{r \times k}$$

$$\text{rank } r \ll \min(d, k)$$

$$W_0 + \Delta W = W_0 + BA$$



Large Model

A 5x5 grid of blue squares, totaling 25 squares.

+

—

LoRA Low-rank Matrices

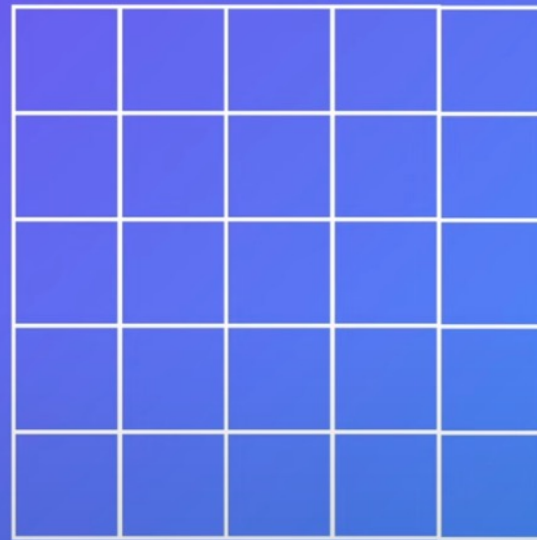



x



=

LoRA Weight Changes




Rank = 1

Increasing Precision by Increasing Rank

LoRA Matrices, Rank 2

x

=

Higher Precision
Weight Changes

Rank = 2

0.15	-0.14	-0.21	0.612
-0.22	0.204	0.308	-0.86
-0.30	-0.16	0.634	0.147
-0.07	-0.2	0.246	0.523

ΔW

Shape: (200, 200)

0.3	-0.14
-0.42	0.201
0.46	0.38
0.5	0.14

B

Shape: (200, 2)

0.1	-0.44	0.04	1.42
-0.92	0.1	1.62	-1.33

A

Shape: (2, 200)

```
def regular_forward_matmul(x, W):
```

```
    h = x @ W
```

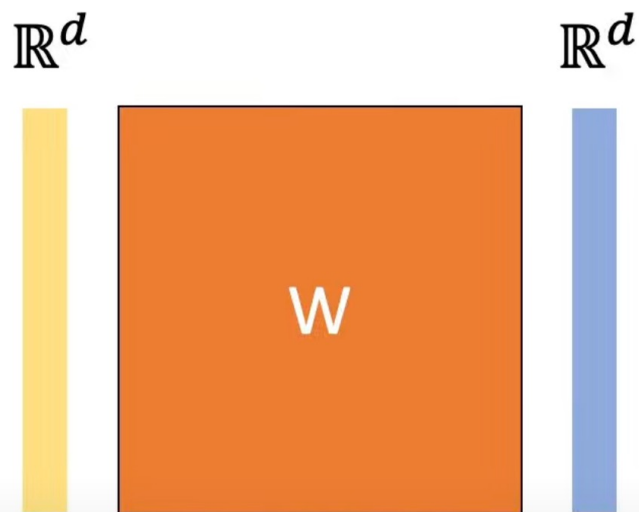
```
    return h
```

```
def lora_forward_matmul(x, W, W_A, W_B):
```

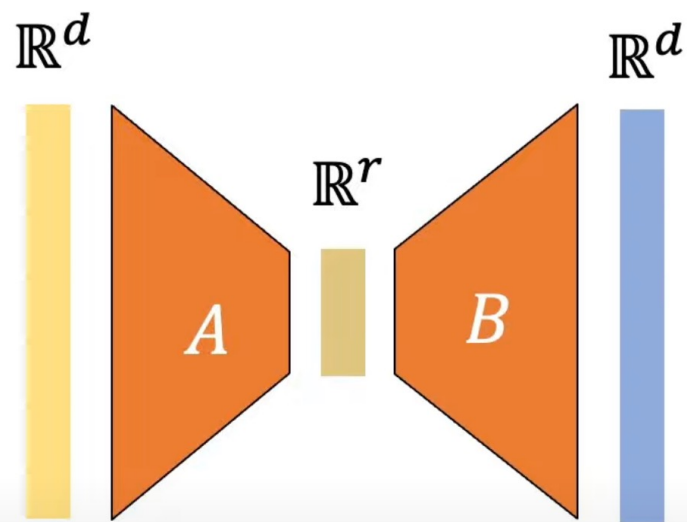
```
    h = x @ W # regular matrix multiplication
```

```
    h += x @ (W_A @ W_B)*alpha # use scaled LoRA weights
```

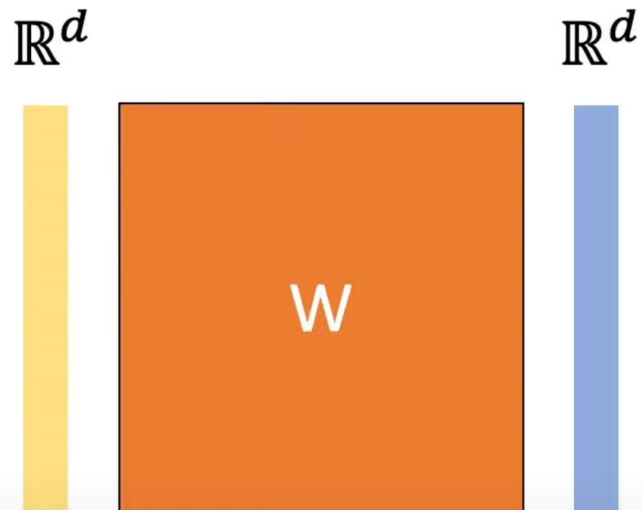
```
    return h
```



Full rank (rank = d)

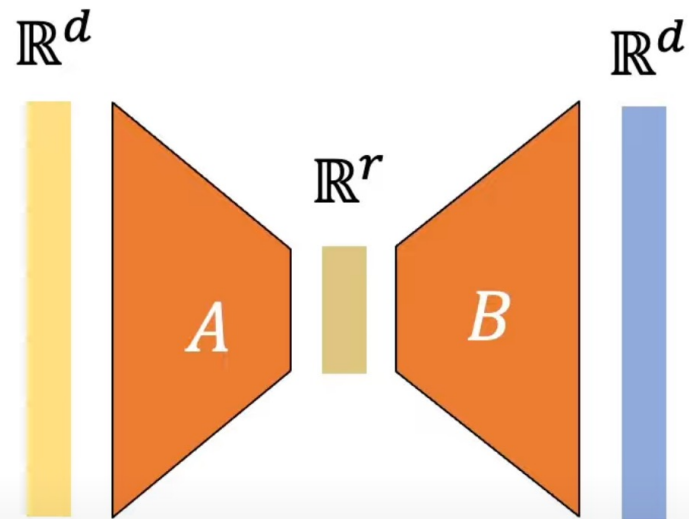


Low rank (rank = $r \ll d$)



Full Rank

$d \times d$ parameters



Low Rank

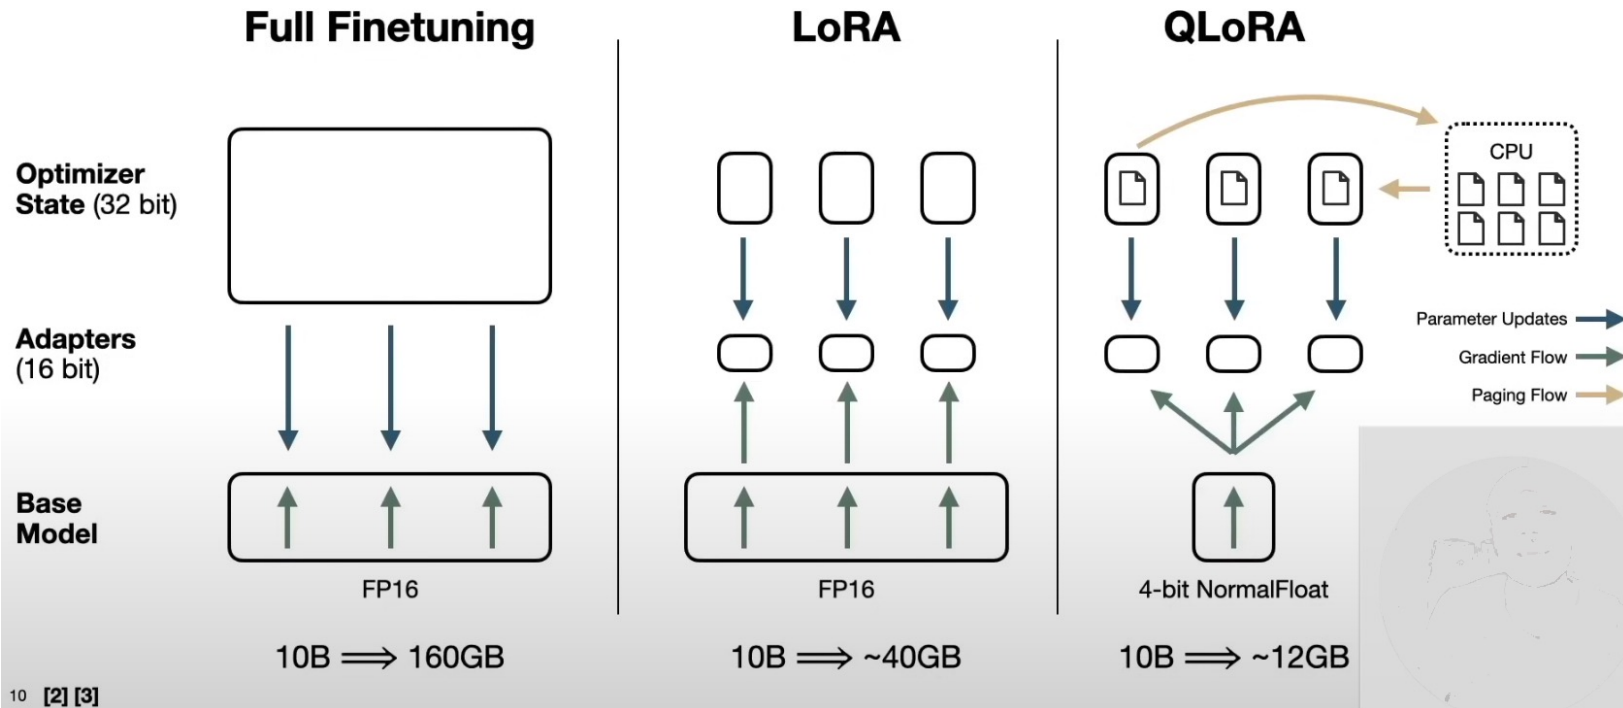
$2 \times d \times r$ parameters

Number of Trainable Parameters

Rank	7B	13B	70B	180B
1	167k	228k	529k	849k
2	334k	456k	1M	2M
8	1M	2M	4M	7M
16	3M	4M	8M	14M
512	86M	117M	270M	434M
1,024	171M	233M	542M	869M
8,192	1.4B	1.8B	4.3B	7.0B

In reality, LLMs are made up of multiple layers of differing sizes. This is a generalization as if the model were a single layer.

Bringing it all together



```

from peft import LoraConfig, get_peft_model

config = LoraConfig(
    r=16, # attention heads
    lora_alpha=32, # alpha scaling
    target_modules=["query_key_value"], # gathered from print(model)
    lora_dropout=0.05,
    bias="none",
    task_type="CAUSAL_LM" # set this for CLM or Seq2Seq
)

model = get_peft_model(model, config)

```

```

print(model)

BloomForCausalLM(
  (transformer): BloomModel(
    (word_embeddings): Embedding(250880, 1024)
    (word_embeddings_layernorm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
    (h): ModuleList(
      (0-23): 24 x BloomBlock(
        (input_layernorm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
        (self_attention): BloomAttention(
          (query_key_value): Linear(in_features=1024, out_features=3072, bias=True)
          (dense): Linear(in_features=1024, out_features=1024, bias=True)
          (attention_dropout): Dropout(p=0.0, inplace=False)
        )
        (post_attention_layernorm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
        (mlp): BloomMLP(
          (dense_h_to_4h): Linear(in_features=1024, out_features=4096, bias=True)
          (gelu_impl): BloomGelu()
          (dense_4h_to_h): Linear(in_features=4096, out_features=1024, bias=True)
        )
      )
    )
    (ln_f): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
  )
  (lm_head): Linear(in_features=1024, out_features=250880, bias=False)
)

```



PEFT

```
from transformers import AutoModelForSeq2SeqLM
from peft import get_peft_config, get_peft_model, LoraConfig, TaskType
model_name_or_path = "bigscience/mt0-large"
tokenizer_name_or_path = "bigscience/mt0-large"

peft_config = LoraConfig(
    task_type=TaskType.SEQ_2_SEQ_LM, inference_mode=False, r=8, lora_alpha=32, lora_dropout=0.1
)

model = AutoModelForSeq2SeqLM.from_pretrained(model_name_or_path)
model = get_peft_model(model, peft_config)
model.print_trainable_parameters()
# output: trainable params: 2359296 || all params: 1231940608 || trainable%: 0.19151053100118282
```